



High Performance Histograms as *Objects*

Henry Schreiner, Hans Dembinski, Jim Pivarski, Shuo Liu



A histogram is best described as an object.

Current Status Quo

`np.histogram`

`plt.hist`



Unbinned data → bins, edge arrays



1D, 2D, ND have different functions/APIs

Hard to plot *results* of `np.histogram`

Generalized histograms:
`scipy.stats.binned_statistic`

And how about binning/manipulating histograms afterwards?

What would a “Histogram” look like?



Axes:
Regular,
Variable,
Category,
...

Accumulators:
Int,
Double,
WeightedSum
Mean,
...

Basics of boost-histogram

Basics:

```
h = bh.Histogram(  
    bh.axis.<Type>(),  
    bh.axis.<Type>(),  
    ...,  
    storage=bh.storage.<Type>())
```

```
h.fill(data, weights=...  
        samples=...)
```

```
# Optional classic style  
bins, *edges = h.to_numpy()
```


← But avoid this, use h directly!

A taste of manipulation

Assume we have a hist “h”:

Plot a histogram: (simple)

```
plt.plot(*h.axes.centers, h, ".")
```



NumPy comparison:

```
bins, edges = h.to_numpy()
centers = (bins[1:] + bins[:-1]) / 2
plt.plot(centers, bins, ".")
```

Compute the density:

```
V = np.prod(h.axes.widths, axis=0)
density = h.view() / h.sum() / V
```

A taste of manipulation

Assume we have a hist “h”:

Manual threaded fill:


```
def fun(d):  
    return h.copy().reset().fill(d)  
  
chunks = np.array_split(data, threads)  
  
with ThreadPoolExecutor(threads) as pool:  
    results = pool.map(fun, chunks)  
  
for res in results:  
    h += res
```

Or just use this:


```
h.fill(data, threads=threads)
```


A taste of manipulation (2)


```
h = bh.Histogram(
    bh.axis.Regular(
        300, 0, 1,
        metadata="x"
    ),
    bh.axis.Regular(
        300, 0, 1,
        metadata="y"
    )
)
```



h



`h[:, :, bh.rebin(6), :, bh.rebin(6)]`



`h[bh.loc(.7):, bh.loc(.5):bh.loc(.8)]`

And boost-histogram is *fast*

Tests on 2.4 GHz 8-Core Intel Core i9

1D, 100 bins, 10,000,000 data points

`bh.numpy.histogram(data, bins, ranges)`

`np.histogram(data, bins, ranges)`

43.1 ms

(41.6 ms in object mode)

74.5 ms

13.8 ms

(13.3 ms in object mode)

Threaded

2D, 100x100 bins, 10,000,000x2 data points

`bh.numpy.histogram2d(data, bins, ranges)`

`np.histogram2d(data, bins, ranges)`

84.7 ms

(77.6 ms in object mode)

874 ms

29.6 ms

(28.7 ms in object mode)

Threaded

A firm foundation



Seamless operability between C++11 and Python, used by SciPy, 1,934 GitHub repos



...one of the most highly regarded and expertly designed C++ library projects in the world.
— Herb Sutter and Andrei Alexandrescu, C++ Coding Standards



Designed by Hans Dembinski, accepted into Boost 1.70, improved every version since.

Boost.Histogram features

- Static / Dynamic storage (can avoid allocator!)
- Static fill becomes 57 lines of vectorized assembly
- Adding, scaling, slicing, rebinning, projections, more
- High performance filling and bin iteration

Customizable:

- Storage
- Allocators
- Accumulators
- Axes
- Axis metadata
- Axis transforms

10

Boost.Histogram author closely involved in boost-histogram!

New bindings: boost-histogram

Design

Flexibility

Performance

Distribution

The bindings
were designed
around these four
key areas

Design

Pickle

Directly supports pickling (v>0)
Optimized for performance
Cloudpickle supported too

Copy

Supports `h.copy()`, `copy(h)`,
and `deepcopy(h)`

Operators

`h1 + h2`
`h * 2.0`

UFuncts

NumPy 1.13 overloads being adopted
for accumulator storages

Axes access

Axes return an enhanced tuple

`h.axes.centers`

`h.axes.widths`

`h.axes.edges`

`h.index(value)`

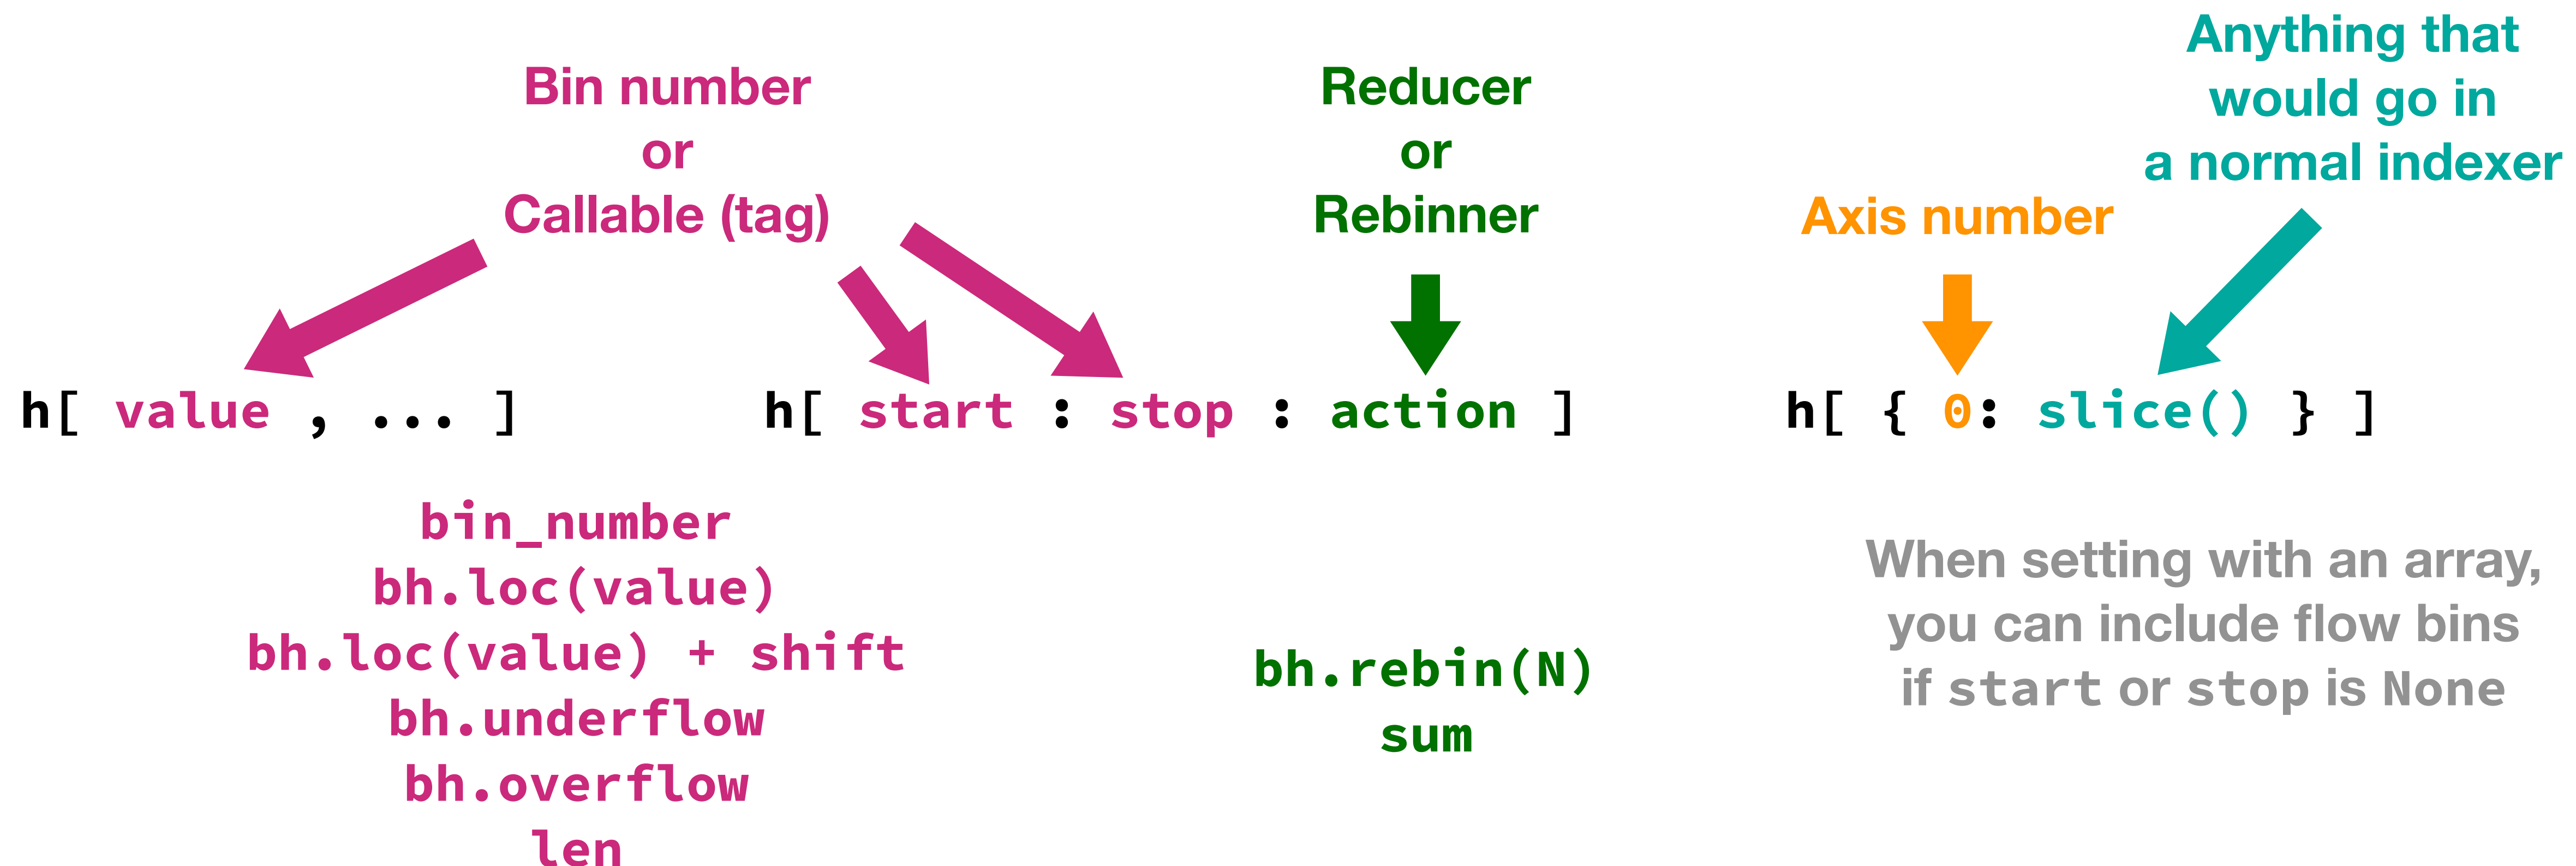
`h.value(index)`

...

Python 2 Statement

1.0 LTS w/ Python 2 support
Python 3 should not suffer

Design: Unified Histogram Indexing











When setting with an array, you can include flow bins if start or stop is None

Any callable works, takes the axes and returns bin number, from `-1` to `len(ax)+1`

Currently only these implemented, UHI spec includes arbitrary actions

Flexibility: Axes Types



Flexibility: Transforms

Regular axes scale better than variable, no sorted lookup.

If the irregularity is functional, you can avoid the lookup!

```
h = bh.Histogram(  
    bh.axis.Regular(  
        10, 1, 10,  
        transform=bh.axis.transform.log  
    )  
)
```

Flexibility: Transforms

Regular axes scale better than variable, no sorted lookup.

If the irregularity is functional, you can avoid the lookup!

You are not just limited to precompiled transforms (sqrt, log, Pow)!

```
import numba
```

```
@numba.cfunc(numba.float64(numba.float64))
```


```
def exp(x):
    return math.exp(x)
```

```
@numba.cfunc(numba.float64(numba.float64))
```


```
def log(x):
    return math.log(x)
```

```
bh.axis.Regular(10, 1, 4,
                transform=bh.axis.transform.Function(log, exp))
```

```
plt.plot(*ho.axes.centers, ho, '.')
```



```
plt.plot(*hl.axes.centers, ho, '.')
```



You can use any ctypes function pointer, or pure Python (slower)

Flexibility: Storages

Boost-histogram ships with 7 storages:


Double() (default)
Fast, flexible, simple

Int64()
Fast, strict, simple

AtomicInt64()
Threadsafe fills

Unlimited()
No overflow guarantee,
Resizes to optimize small ints

Accumulator storages



Weight()
Stores value and variance

Mean()
High accuracy

WeightedMean()
Mean and weight

Views into accumulators are “smart”:

- NumPy record structure
- Property access to real and computed values
- NEP 13 (NumPy 1.13+) UFunc support
- Accessing an element returns an accumulator

```
h.view().value
h.view().variance
np.sum(h.view())
```

**The accumulators (including Sum())
act like 0D histograms, and can be filled!**

Performance

Tests on 2.4 GHz 8-Core Intel Core i9

1D, 100 bins, 10,000,000 data points

2D, 100x100 bins, 10,000,000x2 data points

Setup	Single threaded	X	Multithreaded	X
NumPy 1D	74.5 ± 2.4 ms	1		
BH 1D	41.6 ± 0.7 ms	1.8	13.3 ± 0.2 ms	5.5
BHNP 1D	43.1 ± 0.8 ms	1.7	13.8 ± 0.2 ms	5.4
NumPy 2D	874 ± 22 ms	1		
BH 2D	77.6 ± 0.6 ms	11	28.7 ± 0.7 ms	30
BHNP 2D	85 ± 3 ms	10	29.6 ± 0.5 ms	29

Performance: Single Histogram loop

```

value_ax = bh.axis.Regular(100, -5, 5)
valid_ax = bh.axis.Integer(0, 2,
                           underflow=False,
                           overflow=False)
label_ax = bh.axis.StrCategory([], growth=True)

hist = bh.Histogram(value_ax, valid_ax, label_ax)

hist.fill([-2, 2, 4, 3],
          [True, False, True, True],
          ["a", "b", "a", "b"])
# Just valid data, combine all labels
all_valid = hist[:, bh.loc(True), ::sum]

# 2D histogram of just the "a" label
a_only = hist[..., bh.loc("a")]

```

← Loops over the data just once!

Distribution

Making a distribution:

- Needs to work everywhere
- + Header-only Boost + PyBind11
- C++14

Wheels

Python	2.7	3.5	3.6	3.7	3.8
Manylinux1*	✓	✓	✓	✓	✓
Manylinux2010	✓	✓	✓	✓	✓
macOS 10.9+	✓	✓	✓	✓	✓
Windows	✱	✓	✓	✓	✓

Conda-Forge

Python	2.7	3.5	3.6	3.7	3.8
Linux 64	✱	✓	✓	✓	✓
Linux ppc64le		✓	✓	✓	✓
Linux aarch64		✓	✓	✓	✓
macOS 10.9+	✱	✓	✓	✓	✓
Windows		✓	✓	✓	✓

Solutions:

- git submodules for Boost + Pybind11
- Fully setuptools-based build, optional CMake build
- Custom wheel build system
 - Custom image with GCC 9 (manylinux1)
 - Manylinux2014 arch wheels under investigation
 - Replaced by cibuildwheel recently
 - 4-5 other Scikit-HEP packages now provide wheels
- SDist builds on a modern system
 - PEP 518; C++14 support required

<https://iscinumpy.gitlab.io/categories/azure-devops/>

<https://scikit-hep.org/developer>

Distribution: checks

Pre-commit & GHA:

- Black
- Pre-commit-hooks
- Flake8 with bugbear
- MyPy (basic)
- Check Manifest
- Clang Format (C++)

Test suite:

- PyTest
- PyTest-Benchmark
- Weekly dependency check

Releases:

- Automated by GHA
- setuptools_scm for versioning

See <https://scikit-hep.org/developer>

Several other packages in Scikit-HEP are now adopting recommendations and build tools!



iminuit

Awkward
Array

Particle

Decay
Language

Boost
Histogram

The Scikit-HEP ecosystem

Focused packages working together.



Convert between formats

- NumPy, Physt, and HEP-specific formats



Fill and manipulate histograms

- Core library


histoprint:

Display histograms on the command line

- Newest member of Scikit-HEP
- Can plot up to 5 at a time

Hist:
Friendly histograms for analysis

- GSoC project just started



Extensions for HEP plotting

- Histograms and more



Scikit-HEP Tutorials:

Packages working together

- Powered by JupyterBook Beta

Hist

Python 3.6+ library for users

Testing useful shortcuts (some *may* be upstreamed if popular)

`h[bh.loc(.7):, bh.loc(.5):bh.loc(.8)+1] → h[.7j, .5j:.8j+1]`


`h[bh.loc("hi")] → h["hi"]`

Assign meaning to metadata
name, title, units, etc.

Fill and access by name instead of position
(NamedHist enforces name access)

Direct access to plots
1D, 2D, pull plots, and more

Allowed to have dependencies
such as matplotlib



Early example of pull plot

Summary

Design

Built on a firm foundation
Can handle complex situations with ease

Flexibility

Dozens of axes types x 7 storages
Can be extended in the future

Performance

Fast, efficient filling loops, with threads
Can reduce the number of histograms used

Distribution

Easy to build on a modern system
Wheels for every platform
Conda-forge support as well

Easy to install:

```
pip install boost-histogram  
Or  
conda install boost-histogram -c conda-forge
```

Easy to convert from NumPy:

```
np.histogram* → bh.numpy.histogram*  
Add histogram=bh.Histogram to return object  
Use .to_numpy() to get NumPy tuple back
```

**And more exciting developments
coming this Summer, like Hist!**

A histogram is best described as an object.

Hopefully you now agree with me.

And boost-histogram is a great example of one!

Acknowledgments

Support for this work was provided by the National Science Foundation cooperative agreement OAC-1836650 (IRIS-HEP) and OAC-1450377 (DIANA/HEP).

<https://boost-histogram.readthedocs.io>

<https://github.com/scikit-hep/boost-histogram>

Reference links

<https://root.cern>

<https://scipy.org>

<https://numpy.org>

<https://iris-hep.org>

<https://scikit-hep.org>

<https://www.boost.org>

<https://numba.pydata.org>

<https://pandas.pydata.org>

<https://github.com/scikit-hep/hist>

<https://github.com/pybind/pybind11>

<https://github.com/joerick/cibuildwheel>

<https://github.com/scikit-hep/histoprint>

<https://github.com/scikit-hep/awkward-1.0>