


# Building binary extensions

## with pybind11, scikit-build, and cibuildwheel

SciPy 2022 - Austin



Maintainers track

July 15, 2022

Henry Schreiner  
Princeton University

Joe Rickerby  
Nord Projects

Ralf Grosse-Kunstleve  
Google

Wenzel Jakob  
Ecole Polytechnique Fédérale de Lausanne

Matthieu Darbois  
PyPA

Aaron Gokaslan  
Facebook Research

Jean-Christophe  
Fillion-Robin  
Kitware

Matt McCormick  
Kitware

# Package: *pybind11*

## Header-only pure C++11 CPython/PyPy interface

Trivial to add to a project  
No special build requirements

No dependencies  
No precompile phase  
Not a new language

**Designed for one purpose!**

**Think of it like the missing C++ API for CPython**

# Example of usage

```
#include <pybind11/pybind11.h>
```

Standard include

```
int add(int i, int j) {  
    return i + j;  
}
```

Normal C++ to bind

```
PYBIND11_MODULE(example, m) {  
    m.def("add", &add);  
}
```

Create a Python module

Signature statically inferred

Docs and parameter names optional

```
g++ -shared -fPIC example.cpp $(pipx run pybind11 --includes) -o example$(python3-config --extension-suffix)
```

Complete, working, no-install example (linux version)!

# Many great features

## **Simple enough for tiny projects**

714K code mentions of `pybind11` on GitHub

## **Powerful enough for huge projects**

SciPy, PyTorch, dozens of Google projects

## **NumPy support without NumPy headers**

No need to lock minimum NumPy at build

## **Buffer protocol and array classes**

Includes Eigen support too

## **Most STL containers and features supported**

Including C++17 additions, like `std::variant` (or boost)

## **Trampoline classes and multiple inheritance**

Complex C++ is supported

## **Vectorize methods or functions**

`py::vectorize`, even on a lambda function

## **Supports interop in both directions**

Can even be used to embed Python in C++

## **Powerful object lifetime controls**

`py::keep_alive`, `std::shared_ptr`, and more

## **Complete control of the Python representation**

Special methods, inheritance, pickling, and more

## **Small binaries prioritized**

Perfect for WebAssembly with Pyodide

## **Cross-extension ABI**

One extension can return a type wrapped in another one

# New in 2.8

2.7 was released at last SciPy!

**py::raise\_from**  
Chaining exceptions

**Local exception mapping**  
Custom rules for each module

**Array improvements**  
View and reshape arrays

**Custom \_\_new\_\_ support**  
Fixing old bug

**Python builtins as callbacks**  
Better consistency

**Embedded interpreter improvements**  
argv setting, \_\_file\_\_ set

**py::slice None support**  
Using std::optional

**Rerun CMake with different Python**  
Updates cache if stale

**mingw support added**  
Tested in CI now too

**py::custom\_type\_setup**  
Advanced GC tricks and more

# New in 2.9

## **py::args can be followed**

No longer required to be last

## **More built-in exception chaining**

Using recent py::raise\_from addition

## **py::const\_name replaces py::\_**

More readable, gettext clash removed

## **py::multiple\_inheritance detection**

Only explicitly required if bases are hidden

## **Better support for std::string\_view**

Interoperable with built-in types

## **Codebase formatted with clang-format**

Using clang-format wheel & pre-commit

**Lots of new checks, small fixes, better support for PyPy, CPython 3.11, etc.**

# New in 2.10

## **Removed Python 2.7 & 3.5 support**

Removed 1K+ LoC, simpler, cleaner

## **New docs theme (Furo)**

Cleaner, ToC for pages, dark-mode

## **Removed MSVC 2015 (& limited 2017)**

Hard to find & not very useful with 3.5 gone

## **Support `std::monostate`**

`std::variant` as optional or non-constructible

## **Python 3.11 beta support**

First version working with all changes

## **`py::capsule::set_name`**

Can manipulate later (DLPack)

## **Better exception handling**

Easier without Python 2.7

## **NumPy type enhancements**

More accessors and type number constructor

## **`py::any set` and `py::frozenset`**

Copy to set (like `py::set`)

## **CMake fixes**

Needed for Pyodide (web assembly) & VCPKG

# 2.10.0 releasing today

# Upcoming plans

<https://github.com/pybind/pybind11/wiki/Roadmap>

## **Refactors**

Decoupling unit tests  
Breaking into smaller, IWUY headers

## **Merge the smart holder branch**

Full interoperability with `std::unique_ptr` and `std::shared_ptr`  
Opt-in with `<pybind11/smart_holder.h>`  
Safe passing to C++, including trampolines

## **Optional Precompilation**

Would dramatically speed up compile  
Huge change, so being held back by other work

## **Further *ideas* (not promised yet)**

Better MyPy integration (some minor work done)  
Better Scikit-build integration  
Upstreaming some improvements from nanobind (next slide)



# Maintenance: Release notes

Somewhat interesting release note system

## Pull requests get a template to fill out

```
<!--  
Title (above): please place [branch_name] at the beginning if you are targeting a branch other than master.  
*Do not target stable*. It is recommended to use conventional commit format, see conventionalcommits.org, but  
not required.  
-->  
## Description  
  
<!-- Include relevant issues or PRs here, describe what changed and why -->  
  
## Suggested changelog entry:  
  
<!-- Fill in the below block with the expected RestructuredText entry. Delete if no entry needed;  
      but do not delete header or rst block if an entry is needed! Will be collected via a script. -->  
  
``rst  
  
...  
  
<!-- If the upgrade guide needs updating, note that here too -->
```

# Maintenance: Release notes (2)

GHA bot tags PR's after they get merged

```
name: Labeler

on:
  pull_request_target:
    types: [closed]

jobs:
  label:
    name: Labeler
    runs-on: ubuntu-latest
    steps:
      - uses: actions/labeler@main
        if: github.event.pull_request.merged == true
        with:
          repo-token: ${ secrets.GITHUB_TOKEN }
          configuration-path: .github/labeler_merged.yml
```

You might need  
permissions: contents: ...

+ some labeler config files

# Maintenance: Release notes (3)

Script runner in nox

```
@nox.session(reuse_venv=True)
def make_changelog(session: nox.Session) -> None:
    """
    Inspect the closed issues and make entries for a changelog.
    """

    session.install("ghapi", "rich")
    session.run("python", "tools/make_changelog.py")
```

Controls script dependencies

```
~/g/s/pybind11  skylion007/311-testing $... nox -s make_changelog
```

182ms < Thu Jul 14 22:42:52 2022

```
nox > Running session make_changelog
nox > Re-using existing virtual environment at .nox/make_changelog.
nox > python -m pip install ghapi rich
nox > python tools/make_changelog.py
```

★ optimizes Eigen sparse matrix casting by removing unnecessary temporary.

`#4064 <<https://github.com/pybind/pybind11/pull/4064>>`\_

★ .

`#4060 <<https://github.com/pybind/pybind11/pull/4060>>`\_

★ Ensure proper behavior when garbage collecting classes with dynamic attributes in Python >=3.9.

`#4051 <<https://github.com/pybind/pybind11/pull/4051>>`\_

★ A couple long-standing `PYBIND11\_NAMESPACE` `\_\_attribute\_\_((visibility("hidden")))` inconsistencies are now fixed (affects only unusual environments).

`#4043 <<https://github.com/pybind/pybind11/pull/4043>>`\_

★ Implicit conversion of the literal `0` to `pybind11::handle` is now disabled.

`#4008 <<https://github.com/pybind/pybind11/pull/4008>>`\_

★ New theme for the documentation.

`#3109 <<https://github.com/pybind/pybind11/pull/3109>>`\_

-----  
**Missing:** chore: bump clang-tidy to 13

<https://github.com/pybind/pybind11/pull/3997>

**Template:**

**## Suggested changelog entry:**

```
```rst
```

```
...
```

```
nox > Session make_changelog was successful.
```

```
~/g/s/pybind11  skylion007/311-testing $...
```

2609ms < Thu Jul 14 22:42:58 2022

PRs can be updated, final polish on copy/paste  
Labels manually removed (in bulk)

# Bonus package: nanobind




**C++17+ & Python 3.8+ only**

**Similar API to pybind11**

**Intentionally more limited than pybind11**

**Focus on small, efficient bindings**

**Some ideas can be backported to pybind11**



<https://github.com/wjakob/nanobind>



# Bonus package: build

`pipx run build`

**Builds SDist from source, wheel from SDist**

`pipx run build --sdist --wheel`

**Builds SDist and wheel from source**

`.github/workflows/wheels.yml`

## SDist and wheel builder

### New features this year

**Prettier printout**  
**Better error messages**  
**Python 3.11 support**

### Planned

**Use Flit (better bootstrapping)**  
**Redesign environment support**  
**Threadsafe API**

```
on: [push, pull_request]
```

```
jobs:
```

```
  build_wheels:
```

```
    runs-on: ubuntu-latest
```

```
    steps:
```

```
      - uses: actions/checkout@v4
```

```
      - name: Build wheels
```

```
        run: pipx run build
```

```
      - uses: actions/upload-artifact@v3
```

```
        with:
```

```
          path: ./wheelhouse/*.whl
```

# Package: cibuildwheel

## Multiplatform redistributable wheel builder

**Supports all major CI providers**

GitHub Action provided too!

**Can run locally, on servers, etc.**

**Affiliated (shared maintainer) with manylinux**

**Close collaboration with PyPy devs**

**Joined the PyPA in 2021**

**Used by matplotlib, mypy, scikit-learn, and more**

### **Supports:**

Targeting macOS 10.9+

Apple Silicon cross-compiling 3.8+

All variants of manylinux (including emulation)

musllinux

PyPy 3.7-3.9

Repairing and testing wheels

Reproducible pinned defaults (can unpin)

**2.0 was released at last SciPy!**

**.github/workflows/wheels.yml**

**on:** [push, pull\_request]

**jobs:**

**build\_wheels:**

**runs-on:** \${{ matrix.os }}

**strategy:**

**matrix:**

**os:**

- ubuntu-22.04
- windows-2022
- macos-11

**steps:**

- **uses:** actions/checkout@v4
- **name:** Build wheels  
**uses:** pypa/cibuildwheel@v2.8.0
- **uses:** actions/upload-artifact@v3  
**with:**  
**path:** ./wheelhouse/\*.whl

# Manylinux

**manylinux1**

**RHEL 5 (CentOS)**

Support ended Jan 1, 2022  
Updated till CI breaks

**manylinux2010**

**RHEL 6 (CentOS)**

Support ending Aug 1, 2022

**manylinux2014**

**RHEL 7 (CentOS)**

**Default in cibuildwheel**

**manylinux\_2\_24**

**Debian 9**

Support likely to end

**manylinux\_2\_28**

**RHEL 8 (AlmaLinux)**

**musllinux\_1\_1**

**Alpine 3.12**

**MUSL distributions**



# Cibuildwheel: what's new

## **New in cibuildwheel 2.1-2.4**

Local Windows & MacOS runs

TOML overrides array

manylinux2014 default

musllinux

Environment variable passthrough

Experimental Windows ARM

## **New in cibuildwheel 2.5-2.8**

ABI3 wheel support

Build from SDist

tomllib on Python 3.11 (host)

CPython 3.11 beta support

manylinux\_2\_28 support (RHEL 8)

Linux builds from Windows

Podman support

Support for Pythonless wheels

## **In Development**

New setup-python environment

Python 3.6 removed as host

Python 3.11b4 in progress as of today

# Cibuildwheel: what might be

## **Cross compiling Windows ARM**

Supporting setuptools, hopefully more

## **Cross compiling Linux**

Stuck with missing RH dev toolkit trick

## **Local run improvements**

Skip an environment variable

# cibuildwheel tips

## Build wheels locally

```
pipx run cibuildwheel --platform linux
```

## Use pyproject.toml config

Keep the GitHub Action up-to-date!  
`.github/dependabot.yml`

```
version: 2
updates:
  - package-ecosystem: "github-actions"
    directory: "/"
    schedule:
      interval: "daily"
    ignore:
      - dependency-name: "actions/*"
        update-types:
          - version-update:semver-minor
          - version-update:semver-patch
```

# cibuildwheel tips

## Build wheels locally

```
pipx run cibuildwheel --platform linux
```

## Use pyproject.toml config

## Keep the GitHub Action up-to-date!

`.github/dependabot.yml`

```
version: 2
updates:
  - package-ecosystem: "github-actions"
    directory: "/"
    schedule:
      interval: "daily"
```

**No longer needed,  
Dependabot respects version level!  
(v1 -> v2, not v2.0.0)**

# GHA: building a composite action

GitHub Action's new setup-python@v4.1.0 feature is perfect for composite pipx actions!

**name:** mypkg

**description:** 'Runs a package'

**runs:**

**using:** composite

**steps:**

- **uses:** actions/setup-python@v4

**id:** python

**with:**

**python-version:** "3.7 - 3.10"

**update-environment:** false

- **run:** >

**pipx run**

**--python** '\${{ steps.python.outputs.python-path }}'

**--spec** '\${{ github.action\_path }}'

**mypkg**

**shell:** bash

No side effects!

Used in nox &  
cibuildwheel

} Package you want to run

# Package: scikit-build

## CMake based build backend for Python

**Scikit-build is a CMake-setuptools adaptor from KitWare, the makers of CMake**

First introduced as PyCMake at SciPy 2014 and renamed in 2016

Includes CMake for Python and ninja for Python

**CMake and Ninja for Python**

All manylinux archs • Apple Silicon • cibuildwheel • nox

**Currently designed as a setuptools wrapper**

**But that is changing...**

# Pure Python Packaging

```
# pyproject.toml

[build-system]
requires = ["hatchling"]
build-backend = "hatchling.build"

[project]
name = "example"
version = "0.1.0"
```

See <https://scikit-hep.org/developer>  
Or <https://packaging.python.org>

## Not tied to a single solution!

Adding distutils to the stdlib was a disaster  
Setuptools extending distutils was also a disaster  
But modern standards free us from this!

## Many pure-Python backends today!

Hatch, PDM, Flit, and more!  
Even modern setuptools supports this

## But what if you want to add compiled extensions?

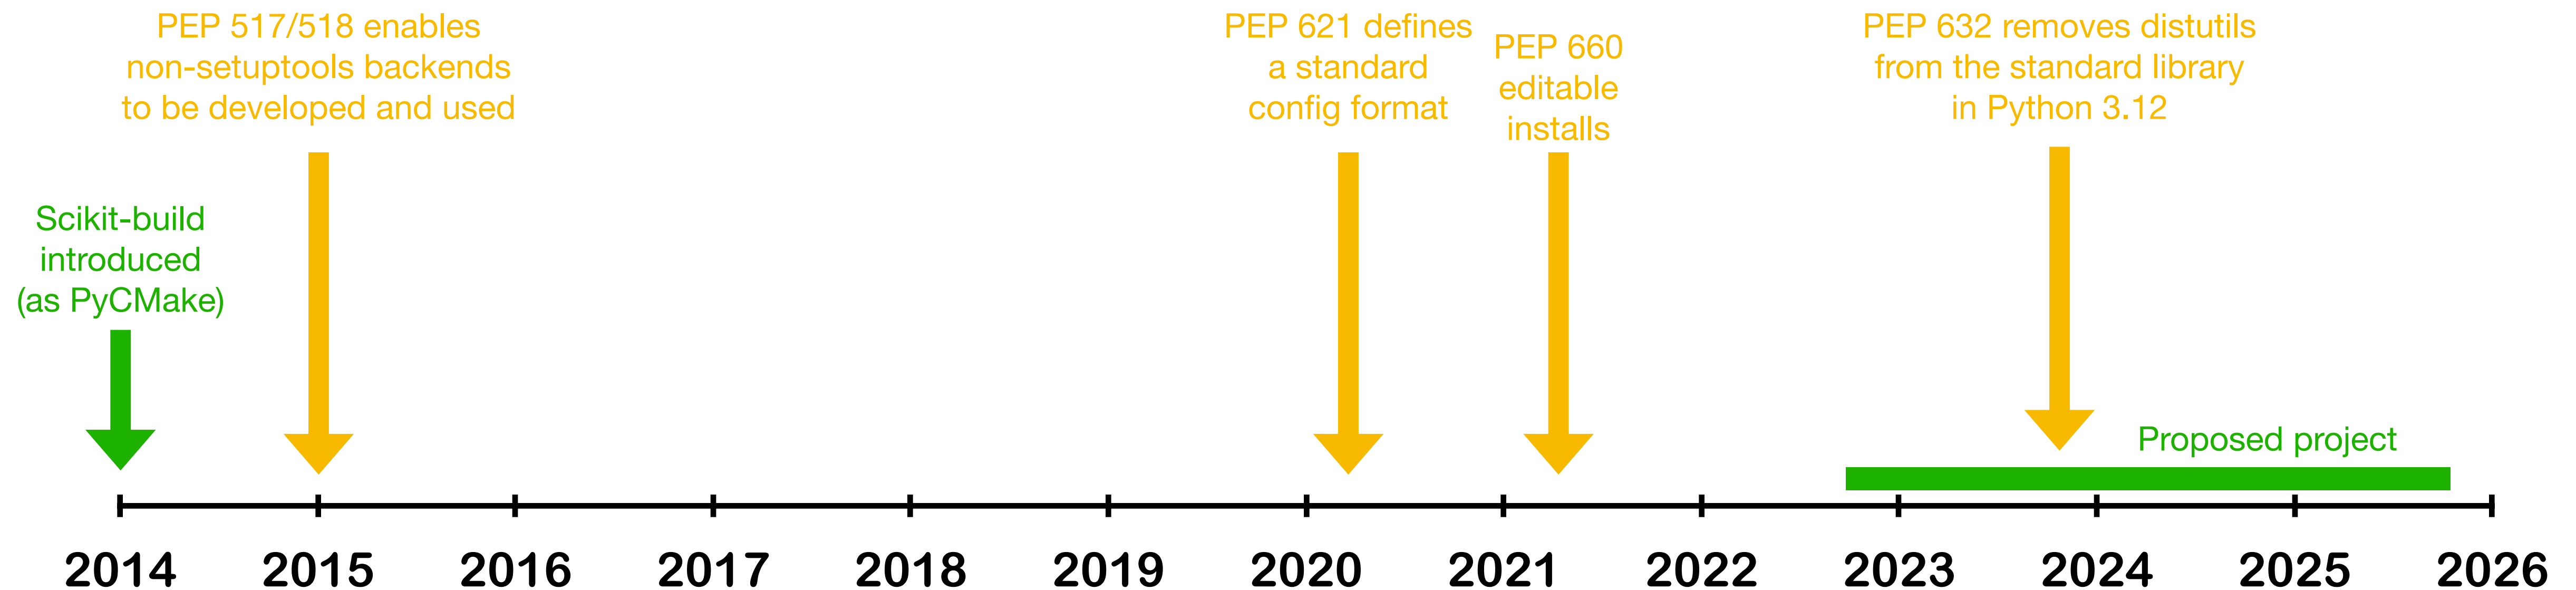
Fall back 8+ years into distutils/setuptools hacks

# Packaging timeline

Before 2015, the only method to distribute packages was distutils/setuptools

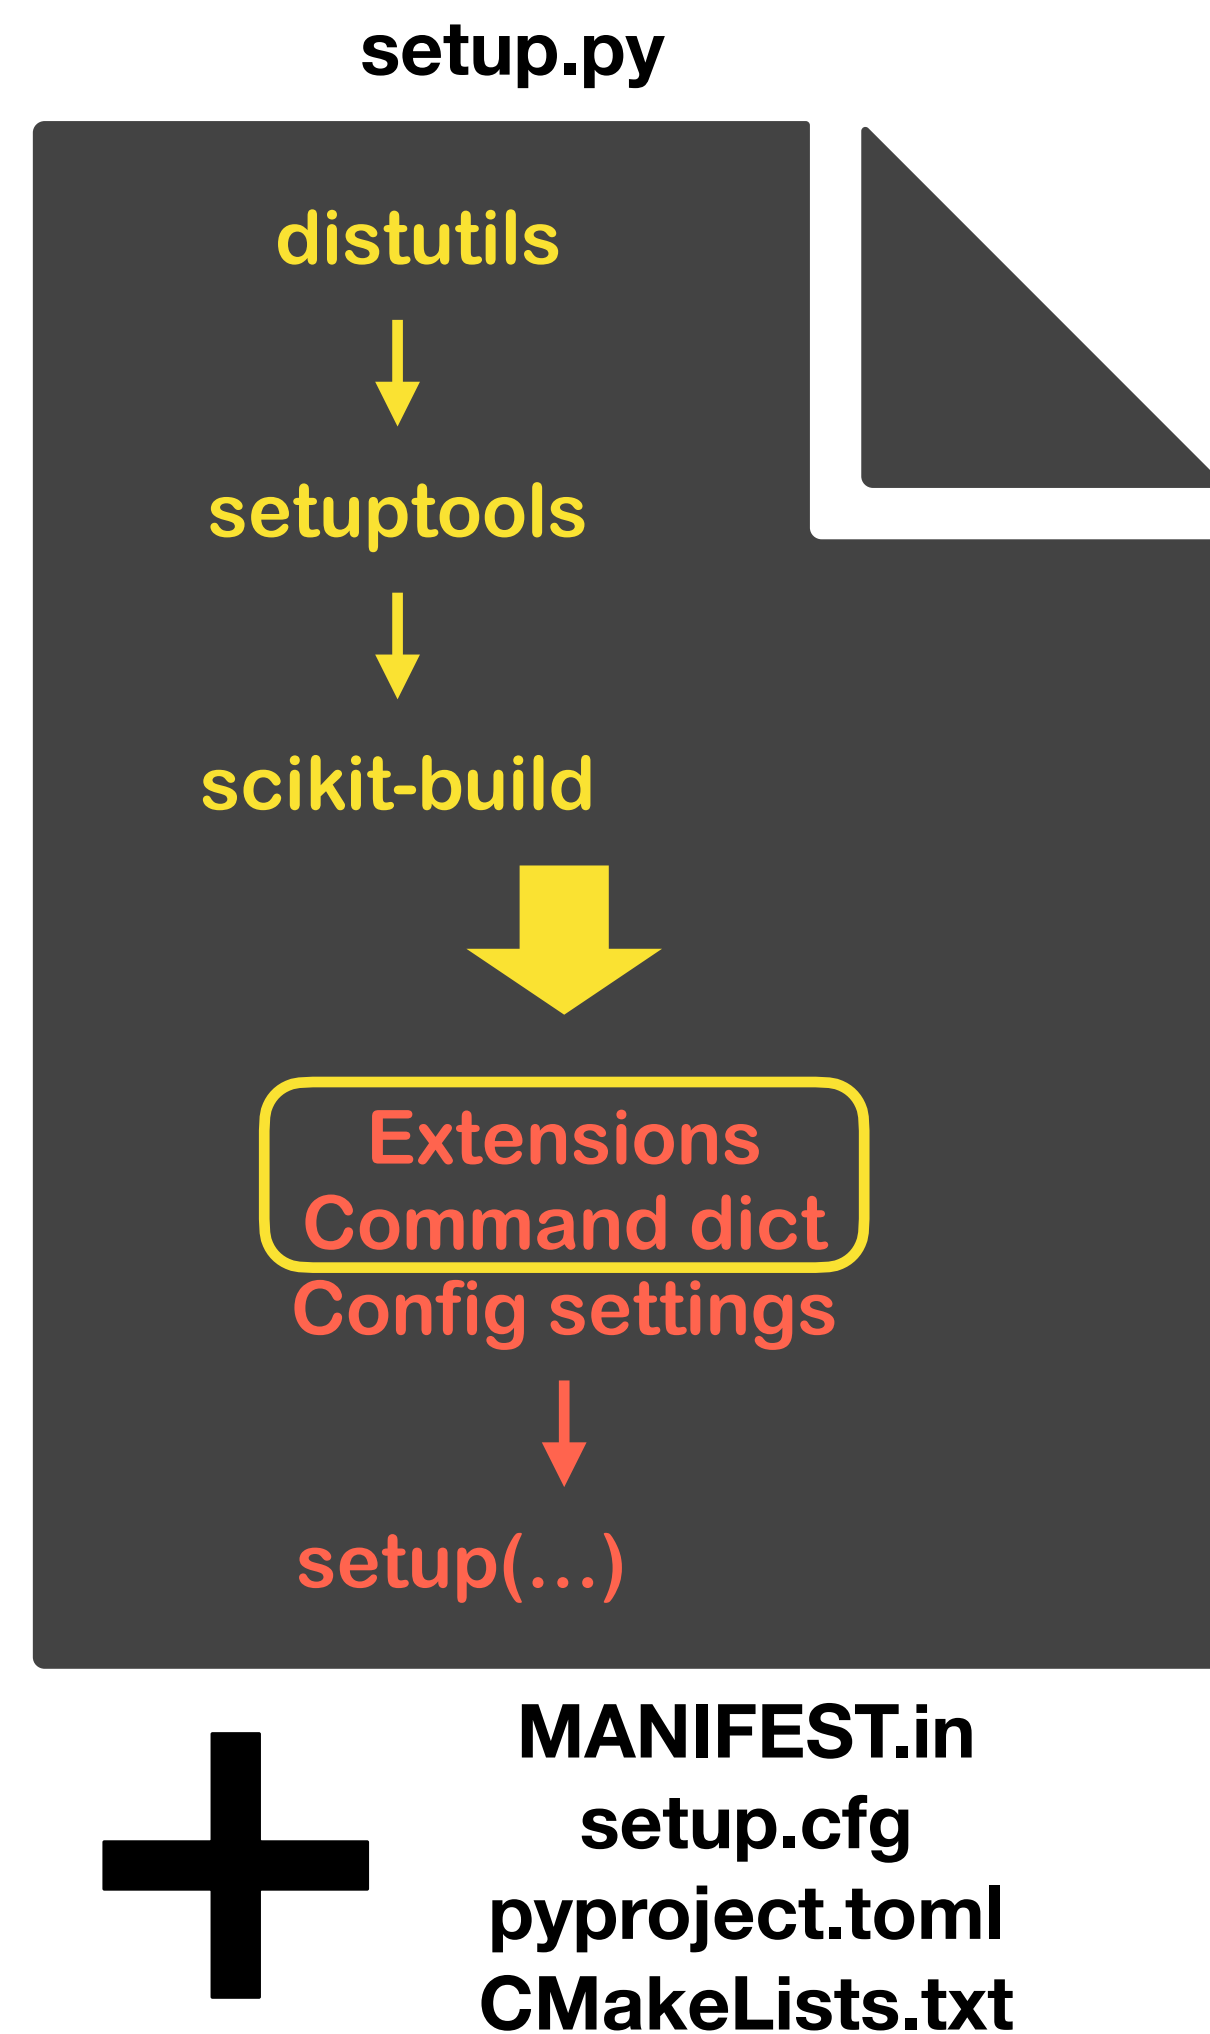
Python Enhancement Proposals (PEPs) drive Python development - and provide standards for packaging now too!

Good pure Python backends exist now (hatchling) - scikit-build could be a great compiled / CMake one!





# Scikit-build (current design)



## Raw setuptools for C++?

No C++ standard support

No parallel compiles

No common libraries

No custom libraries

Hacks for cross compiling

*NumPy alone needs 13K lines to support this!*

## Scikit-build wraps CMake!

Best in class feature support

~60% of all C++ projects

Very popular in the sciences (incl. HEP)

Common and custom library support

## Problems with setuptools bridge

Inherits problems from setuptools

Deep dependencies on internals

Poor interactions with new standards

Deep nesting is hard to work with / debug

Internals poorly documented

Not composable with other plugins


# Scikit-build (new design)

pyproject.toml

build-system  
require  
scikit-build-core

project configuration  
Name, etc.

tool.scikit-build  
Optional custom config



CMakeLists.txt

## Design considerations:

- Good testing suite
- Excellent error messages
- Public API

## Planned other interfaces:

- Extension-based setuptools wrapper
- Can be used to support classic “scikit-build”
- Plugin mode (see next page)

## CMake code:

- Redesign with modern “FindPython” support
- Provide extension discovery system

*You could use pybind11  
from pip in CMake!*

(Scikit-build-core may be a placeholder name)

# Extensionlib (w/ PEP?)

(Rough example of idea)

```
[build-system]
```

```
requires = ["hatchling", "scikit-build-core", "mypyc"]  
build-backend = "hatchling.build"
```

```
[project]
```

```
name = "example"  
version = "0.1.0"
```

```
[[build-system.extensions]]
```

```
build-backend = "scikit_build_core.extension"  
src = "src1/CMakeLists.txt"
```

```
[[build-system.extensions]]
```

```
build-backend = "scikit_build_core.extension"  
src = "src2/CMakeLists.txt"
```

```
[[build-system.extensions]]
```

```
build-backend = "mypyc.extension"  
src = "src/*.py"
```

Since the proposal, “extensionlib” and/or a PEP has been discussed at PyCon 2022. A PoC lib by @ofek exists. Stay tuned!

Provides a common API for build-backends to extension creation tools

We could be a plugin for all builders via a single interface!

# NSF Proposal

## Year 1 plans:

- Introduce scikit-build-core: modern PEP 517 backend
- Support PEP 621 configuration
- Support use as plugin (possibly via extesionlib)
- Tighter CMake integration (config from PyPI packages)
- Distutils-free code ready for Python 3.12

## Year 2 plans:

- Convert projects to Scikit-build

## Year 3 plans:

- Website, tutorials, outreach

<https://iscinumpy.dev/post/scikit-build-proposal/>

**On Wednesday, NSF 2209877 was awarded!**

Updates can touch many repos!

**Scikit-build family:**  
scikit-build, cmake, ninja,  
moderncmakedomain,  
scikit-build-examples

**PyPA family:**  
cibuildwheel,  
build

**pybind family:**  
pybind11,  
python\_example,  
cmake\_example,  
scikit\_build\_example

**scikit-hep family:**  
Developer pages, cookie

# Scikit-build, since last time

## **CMake**

Current release (3.23) will likely drop manylinux1 wheels

## **Scikit-build 0.13**

ARM on Windows support

MSVC 2022 support

Target link libraries support via keyword

## **Scikit-build 0.14**

Install target

## **Scikit-build 0.15**

Cygwin support

(Limited) FindPython support

Final Python 2.7 & 3.5 release

Final MSVC 2015 release

*+ lots of fixes  
& minor improvements*

# Questions?

To get in touch, currently we are using GitHub Discussions on all three packages.

Special thank you to Google for sponsoring our expanded CI, and for Ralf Grosse-Kunstleve's time.

Henry Schreiner supported by IRIS-HEP and the NSF under Cooperative Agreement OAC-1836650



<https://iscinumpy.dev>

<https://scikit-hep.org>  
<https://iris-hep.org>



## C++ & Python

[pybind11](#) ([python\\_example](#), [cmake\\_example](#), [scikit\\_build\\_example](#)) • [Conda-Forge ROOT](#)

## Building Python Packages

[cibuildwheel](#) • [build](#) • [scikit-build](#) ([cmake](#), [ninja](#), [sample-projects](#)) • [Scikit-HEP/cookie](#)

## Scikit-HEP: Histograms

[boost-histogram](#) • [Hist](#) • [UHI](#) • [uproot-browser](#)

## Scikit-HEP: Other

[Vector](#) • [Particle](#) • [DecayLanguage](#) • [repo-review](#)

## Other C++

[CLI11](#) • [GooFit](#)

## Other Ruby

[Jekyll-Indico](#)

## Other Python

[Plumbum](#) • [POVM](#) • [PyTest GHA annotate-failures](#)

## My books and workshops

[Modern CMake](#) • [CMake Workshop](#)

[Computational Physics Class](#)

Python [CPU](#), [GPU](#), [Compiled](#) minicourses

[Level Up Your Python](#)





# Scikit-HEP Developer Pages

Bonus slides (from Scikit-HEP talk yesterday)



# Scikit-HEP developer pages

Scikit-HEP was growing quickly  
How to keep quality high across all packages?  
[scikit-hep.org/developer/](https://scikit-hep.org/developer/)

## Topics

Packaging (classic & simple)  
Style guide (pre-commit)  
Development setup  
pytest  
Static typing (MyPy)  
GitHub Actions (3 pages)  
Nox

## Utilities

Cookie-cutter  
Repo-review

Maintained by nox  
& GitHub Actions

Guided our other packages, like Awkward

Universally useful advice for maintaining packages!



# Scikit-HEP project – welcome!

The Scikit-HEP project is a community-driven and community-oriented project with the aim of providing Particle Physics at large with an ecosystem for data analysis in Python. [Read more -](#)

New users can start with our [user pages](#). See our [developer pages](#) for information on developing Python packages.

[NEWS](#) • [TUTORIAL](#) • [RESOURCES](#) • [CITE US](#) • [GET IN TOUCH](#)

## Basics:

**Awkward  
Array**

Manipulate JSON-like data with NumPy-like idioms.

**hepunits**

Units and constants in the HEP system of units.

**VECTOR**

Manipulate Lorentz, 3D, and 2D vectors in NumPy, Numba, or Aw

Home

Project news

Packages



User information



Developer information



Who uses Scikit-HEP?



About



# Style

Pre-commit hooks

Black

Check-Manifest (setuptools only)

Type checking

PyCln

Flake8

YesQA

isort

PyUpgrade

Setup.cfg format (setuptools only)

Spelling

PyGrep hooks

Prettier

Clang-format (C++ only)

Shellcheck (shell scripts only)

PyLint (noisy)

# Modern (simple) packaging

pyproject.toml

```
[build-system]
requires = ["hatchling"]
build-backend = "hatchling.build"

[project]
name = "package"
version = "0.1.2"
```

No setup.py, setup.cfg, or MANIFEST.in required!

# Cookiecutter

```
pipx run cookiecutter gh:scikit-hep/cookie
```

11 backends to pick from  
Generation tested by nox  
In sync with the developer pages

Setuptools  
Setuptools PEP 621  
Flit Hatch PDM  
Poetry

Scikit-build  
Setuptools C++  
Maturin (Rust)

+ more!



[Home](#)[Project news](#)[Packages](#)[User information](#)[Developer information](#)[Intro to development](#)[Testing with pytest](#)[Simple Python packaging](#)[Packaging](#)[Style](#)[Static type checking](#)[GHA: GitHub Actions intro](#)[GHA: Pure Python wheels](#)[GHA: Binary wheels](#)[Badges](#)

# Scikit-HEP Repo Review (beta)

You can check the style of a GitHub repository below. Enter any repository, such as `scikit-hep/hist`, and the branch you want to check, such as `main` (it must exist). This will produce a list of results - green checkmarks mean this rule is followed, red errors mean the rule is not. A yellow warning sign means that the check was skipped because a previous required check failed. Some checks will fail, that's okay - the goal is bring all possible issues to your attention, not to force compliance with arbitrary checks.

You can also run [this tool](#) locally:

```
pipx run 'scikit-hep-repo-review[cli]' <path to repo>
```

Org/Repo


`pypa/build`

e.g. scikit-hep/hist

Branch

`main`

e.g. main



Results for pypa/build@main

general



PY001: Has a pyproject.toml



PY002: Has a README.(md|rst) file



PY003: Has a LICENSE\* file



PY004: Has docs folder

**Runs in WebAssembly locally!  
(via Pyodide)**